

dl-1

April 11, 2025

```
[7]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[8]: BostonTrain = pd.read_csv("boston_test.csv")
```

```
[9]: BostonTrain.head()
```

```
[9]:
```

	ID	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	
1	6	0.02985	0.0	2.18	0	0.458	6.430	58.7	6.0622	3	222	
2	8	0.14455	12.5	7.87	0	0.524	6.172	96.1	5.9505	5	311	
3	9	0.21124	12.5	7.87	0	0.524	5.631	100.0	6.0821	5	311	
4	10	0.17004	12.5	7.87	0	0.524	6.004	85.9	6.5921	5	311	

	ptratio	black	lstat
0	17.8	392.83	4.03
1	18.7	394.12	5.21
2	15.2	396.90	19.15
3	15.2	386.63	29.93
4	15.2	386.71	17.10

```
[10]: BostonTrain.info()
BostonTrain.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 173 entries, 0 to 172
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   ID           173 non-null    int64
1   crim        173 non-null    float64
2   zn          173 non-null    float64
3   indus       173 non-null    float64
4   chas        173 non-null    int64
5   nox         173 non-null    float64
```

```

6   rm      173 non-null   float64
7   age     173 non-null   float64
8   dis     173 non-null   float64
9   rad     173 non-null   int64
10  tax     173 non-null   int64
11  ptratio 173 non-null   float64
12  black   173 non-null   float64
13  lstat   173 non-null   float64
dtypes: float64(10), int64(4)
memory usage: 19.1 KB

```

```

[10]:
      count  ID      crim      zn      indus      chas      nox  \
count  173.000000  173.000000  173.000000  173.000000  173.000000  173.000000
mean    258.404624    4.100862   12.661850   10.835145    0.086705    0.549981
std     143.289788   10.607761   24.536277    6.596488    0.282219    0.117826
min       3.000000    0.013810    0.000000    0.460000    0.000000    0.392000
25%     136.000000    0.082210    0.000000    5.320000    0.000000    0.447000
50%     268.000000    0.251990    0.000000    8.560000    0.000000    0.538000
75%     381.000000    3.673670   20.000000   18.100000    0.000000    0.624000
max     505.000000   88.976200   95.000000   27.740000    1.000000    0.871000

      count  rm      age      dis      rad      tax      ptratio  \
count  173.000000  173.000000  173.000000  173.000000  173.000000  173.000000
mean     6.321237   69.245665   3.958865    9.387283  406.231214   18.469942
std      0.700621   28.248244   2.324131    8.662621  164.480626    2.196196
min      4.138000    2.900000   1.178100    1.000000  187.000000   12.600000
25%      5.895000   42.800000   2.010700    4.000000  279.000000   17.000000
50%      6.223000   79.200000   3.421100    5.000000  330.000000   19.100000
75%      6.674000   94.600000   5.400700   24.000000  666.000000   20.200000
max      8.780000  100.000000  12.126500   24.000000  711.000000   22.000000

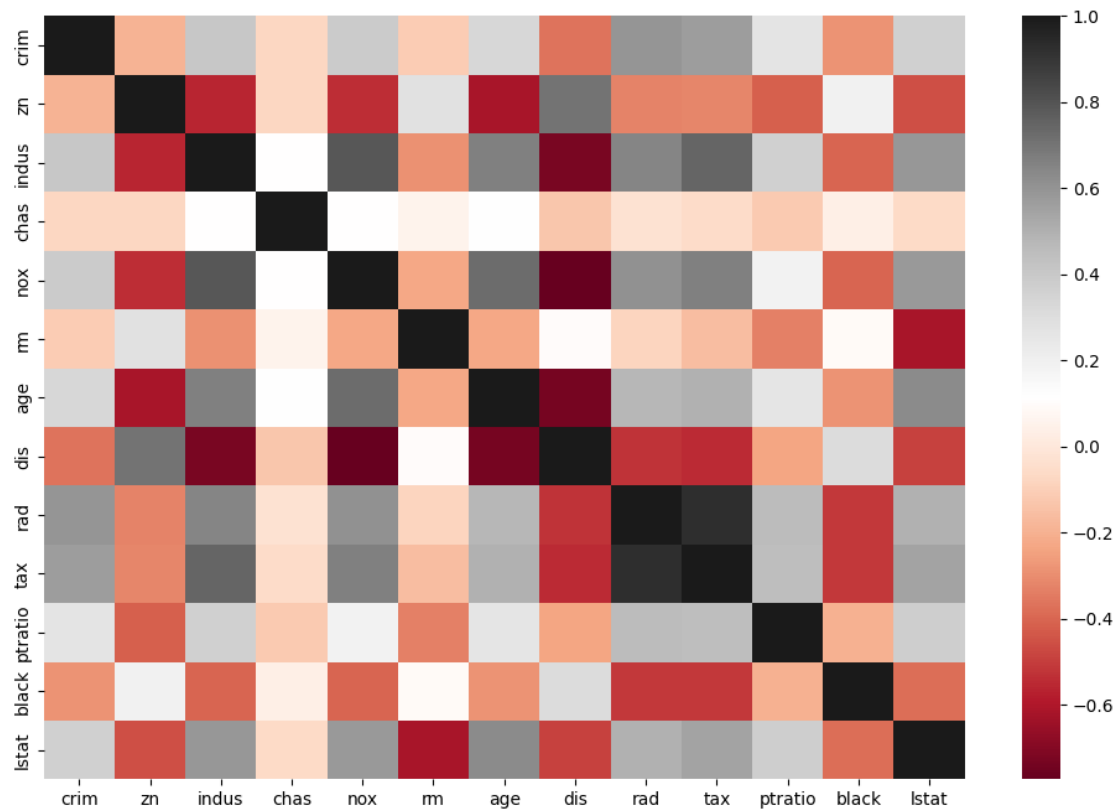
      count  black      lstat
count  173.000000  173.000000
mean    351.299711   12.917977
std     99.781464    7.293408
min      0.320000    1.920000
25%    371.720000    6.870000
50%    390.070000   12.120000
75%    396.060000   17.210000
max    396.900000   34.370000

```

```
[11]: BostonTrain.drop('ID', axis = 1, inplace=True)
```


```
[13]: plt.subplots(figsize=(12,8))
      sns.heatmap(BostonTrain.corr(), cmap = 'RdGy')
```

```
[13]: <Axes: >
```



```
[15]: sns.pairplot(BostonTrain, vars = ['rm', 'zn', 'black', 'dis', 'chas'])
```

```
[15]: <seaborn.axisgrid.PairGrid at 0x7bab0a107150>
```



```
[16]: X = BostonTrain[['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis',
    ↪ 'rad', 'tax',
    'ptratio', 'black', 'lstat']]
y = BostonTrain['rad']
```

```
[17]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
[18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4)
```


```
[19]: lm = LinearRegression()
lm.fit(X_train, y_train)
```

```
[19]: LinearRegression()
```

```
[20]: predictions = lm.predict(X_test)
```

```
[21]: plt.scatter(y_test, predictions)
plt.xlabel('Y Test')
plt.ylabel('Predicted Y')
```

```
[21]: Text(0, 0.5, 'Predicted Y')
```



```
[22]: from sklearn import metrics

print('MAE:', metrics.mean_absolute_error(y_test, predictions))
print('MSE:', metrics.mean_squared_error(y_test, predictions))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

```
MAE: 1.0705722023171153e-14
MSE: 1.5758693959949422e-28
RMSE: 1.2553363676700131e-14
```

```
[23]: sns.distplot((y_test-predictions),bins=50);
```


<ipython-input-23-5f2bc21c0ef7>:1: UserWarning:

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot((y_test-predictions),bins=50);
```



```
[24]: coefficients = pd.DataFrame(lm.coef_,X.columns)
      coefficients.columns = ['coefficients']
      coefficients
```

```
[24]:      coefficients
      crim    -8.635243e-17
      zn      1.804112e-16
      indus  -1.567756e-16
```

chas	1.141646e-14
nox	2.328294e-15
rm	-4.165776e-16
age	-6.071532e-18
dis	-4.302656e-16
rad	1.000000e+00
tax	1.249001e-16
ptratio	2.592869e-16
black	1.144917e-16
lstat	9.454243e-17

[]: